

November 1, 2023

1 Logistic Regression

Exp no.: 8

Aim: Logistic Regression

```
[1]: #Name:Vedant Wankhade  
#Roll no.:74  
#Sec:B  
#Aim:SVM Classifier  
#Year:3rd Year
```

```
[2]: import pandas as pd  
import os  
import matplotlib.pyplot as plt  
import numpy as np  
import seaborn as sns  
from sklearn.model_selection import train_test_split  
import warnings  
warnings.filterwarnings('ignore')
```

```
[3]: os.getcwd()
```

```
[3]: 'C:\\\\Users\\hp\\Downloads'
```

```
[4]: os.chdir('C:\\\\Users\\HP\\Desktop')
```

```
[5]: df=pd.read_csv('framingham.csv')
```

```
[6]: df.head()
```

```
[6]:   male  age  education  currentSmoker  cigsPerDay  BPMeds  prevalentStroke  \  
0     1   39         4.0              0          0.0      0.0              0  
1     0   46         2.0              0          0.0      0.0              0  
2     1   48         1.0              1         20.0      0.0              0  
3     0   61         3.0              1         30.0      0.0              0  
4     0   46         3.0              1         23.0      0.0              0
```

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
0	0	0	195.0	106.0	70.0	26.97	80.0	77.0	
1	0	0	250.0	121.0	81.0	28.73	95.0	76.0	
2	0	0	245.0	127.5	80.0	25.34	75.0	70.0	
3	1	0	225.0	150.0	95.0	28.58	65.0	103.0	
4	0	0	285.0	130.0	84.0	23.10	85.0	85.0	

	TenYearCHD
0	0
1	0
2	0
3	1
4	0

```
[7]: df.tail()
```

```
[7]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
4233     1   50         1.0              1         1.0      0.0
4234     1   51         3.0              1        43.0      0.0
4235     0   48         2.0              1        20.0     NaN
4236     0   44         1.0              1        15.0      0.0
4237     0   52         2.0              0         0.0      0.0
```

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
4233	0	1	0	313.0	179.0	92.0	25.97	
4234	0	0	0	207.0	126.5	80.0	19.71	
4235	0	0	0	248.0	131.0	72.0	22.00	
4236	0	0	0	210.0	126.5	87.0	19.16	
4237	0	0	0	269.0	133.5	83.0	21.47	

	heartRate	glucose	TenYearCHD
4233	66.0	86.0	1
4234	65.0	68.0	0
4235	84.0	86.0	0
4236	86.0	NaN	0
4237	80.0	107.0	0

```
[8]: df.info
```

```
[8]: <bound method DataFrame.info of      male  age  education  currentSmoker
cigsPerDay  BPMeds  \
0         1   39         4.0              0         0.0      0.0
1         0   46         2.0              0         0.0      0.0
2         1   48         1.0              1        20.0      0.0
3         0   61         3.0              1        30.0      0.0
4         0   46         3.0              1        23.0      0.0
...     ...  ...         ...              ...     ...
...     ...  ...         ...              ...     ...
```

4233	1	50	1.0	1	1.0	0.0
4234	1	51	3.0	1	43.0	0.0
4235	0	48	2.0	1	20.0	NaN
4236	0	44	1.0	1	15.0	0.0
4237	0	52	2.0	0	0.0	0.0

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
0	0	0	0	195.0	106.0	70.0	26.97	
1	0	0	0	250.0	121.0	81.0	28.73	
2	0	0	0	245.0	127.5	80.0	25.34	
3	0	1	0	225.0	150.0	95.0	28.58	
4	0	0	0	285.0	130.0	84.0	23.10	
...		
4233	0	1	0	313.0	179.0	92.0	25.97	
4234	0	0	0	207.0	126.5	80.0	19.71	
4235	0	0	0	248.0	131.0	72.0	22.00	
4236	0	0	0	210.0	126.5	87.0	19.16	
4237	0	0	0	269.0	133.5	83.0	21.47	

	heartRate	glucose	TenYearCHD
0	80.0	77.0	0
1	95.0	76.0	0
2	75.0	70.0	0
3	65.0	103.0	1
4	85.0	85.0	0
...
4233	66.0	86.0	1
4234	65.0	68.0	0
4235	84.0	86.0	0
4236	86.0	NaN	0
4237	80.0	107.0	0

[4238 rows x 16 columns]>

```
[9]: df.describe()
```

```
[9]:
```

	male	age	education	currentSmoker	cigsPerDay	\
count	4238.000000	4238.000000	4133.000000	4238.000000	4209.000000	
mean	0.429212	49.584946	1.978950	0.494101	9.003089	
std	0.495022	8.572160	1.019791	0.500024	11.920094	
min	0.000000	32.000000	1.000000	0.000000	0.000000	
25%	0.000000	42.000000	1.000000	0.000000	0.000000	
50%	0.000000	49.000000	2.000000	0.000000	0.000000	
75%	1.000000	56.000000	3.000000	1.000000	20.000000	
max	1.000000	70.000000	4.000000	1.000000	70.000000	

	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	\
--	--------	-----------------	--------------	----------	---------	---

count	4185.000000	4238.000000	4238.000000	4238.000000	4188.000000
mean	0.029630	0.005899	0.310524	0.025720	236.721585
std	0.169584	0.076587	0.462763	0.158316	44.590334
min	0.000000	0.000000	0.000000	0.000000	107.000000
25%	0.000000	0.000000	0.000000	0.000000	206.000000
50%	0.000000	0.000000	0.000000	0.000000	234.000000
75%	0.000000	0.000000	1.000000	0.000000	263.000000
max	1.000000	1.000000	1.000000	1.000000	696.000000

	sysBP	diaBP	BMI	heartRate	glucose \
count	4238.000000	4238.000000	4219.000000	4237.000000	3850.000000
mean	132.352407	82.893464	25.802008	75.878924	81.966753
std	22.038097	11.910850	4.080111	12.026596	23.959998
min	83.500000	48.000000	15.540000	44.000000	40.000000
25%	117.000000	75.000000	23.070000	68.000000	71.000000
50%	128.000000	82.000000	25.400000	75.000000	78.000000
75%	144.000000	89.875000	28.040000	83.000000	87.000000
max	295.000000	142.500000	56.800000	143.000000	394.000000

	TenYearCHD
count	4238.000000
mean	0.151958
std	0.359023
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
[10]: df.isna().sum()
```

```
[10]: male          0
      age           0
      education    105
      currentSmoker  0
      cigsPerDay    29
      BPMeds        53
      prevalentStroke  0
      prevalentHyp   0
      diabetes       0
      totChol        50
      sysBP          0
      diaBP          0
      BMI            19
      heartRate       1
      glucose        388
      TenYearCHD      0
```

dtype: int64

```
[11]: df['glucose'].fillna(value = df['glucose'].mean(),inplace=True)
```

```
[12]: df['education'].fillna(value = df['education'].mean(),inplace=True)
```

```
[13]: df['heartRate'].fillna(value = df['heartRate'].mean(),inplace=True)
```

```
[14]: df['BMI'].fillna(value = df['BMI'].mean(),inplace=True)
```

```
df['cigsPerDay'].fillna(value = df['cigsPerDay'].mean(),inplace=True)
```

```
[15]: df['totChol'].fillna(value = df['totChol'].mean(),inplace=True)
```

```
[16]: df['BPMeds'].fillna(value = df['BPMeds'].mean(),inplace=True)
```

```
[17]: df.isna().sum()
```

```
[17]: male                0
      age                0
      education          0
      currentSmoker      0
      cigsPerDay          29
      BPMeds             0
      prevalentStroke     0
      prevalentHyp        0
      diabetes           0
      totChol            0
      sysBP              0
      diaBP              0
      BMI                0
      heartRate          0
      glucose            0
      TenYearCHD         0
      dtype: int64
```

```
[18]: df.isna().sum()
```

```
[18]: male                0
      age                0
      education          0
      currentSmoker      0
      cigsPerDay          29
      BPMeds             0
      prevalentStroke     0
      prevalentHyp        0
      diabetes           0
```

```

totChol      0
sysBP        0
diaBP        0
BMI          0
heartRate    0
glucose      0
TenYearCHD   0
dtype: int64

```

```
[19]: #Splitting the dependent and independent variables.
```

```

x = df.drop("TenYearCHD",axis=1)
y = df['TenYearCHD']

```

```
[20]: x #checking the features
```

```

[20]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
0         1   39         4.0              0          0.0  0.00000
1         0   46         2.0              0          0.0  0.00000
2         1   48         1.0              1         20.0  0.00000
3         0   61         3.0              1         30.0  0.00000
4         0   46         3.0              1         23.0  0.00000
...
4233      1   50         1.0              1          1.0  0.00000
4234      1   51         3.0              1         43.0  0.00000
4235      0   48         2.0              1         20.0  0.02963
4236      0   44         1.0              1         15.0  0.00000
4237      0   52         2.0              0          0.0  0.00000

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
0                   0              0         0    195.0  106.0   70.0  26.97
1                   0              0         0    250.0  121.0   81.0  28.73
2                   0              0         0    245.0  127.5   80.0  25.34
3                   0              1         0    225.0  150.0   95.0  28.58
4                   0              0         0    285.0  130.0   84.0  23.10
...
4233                  0              1         0    313.0  179.0   92.0  25.97
4234                  0              0         0    207.0  126.5   80.0  19.71
4235                  0              0         0    248.0  131.0   72.0  22.00
4236                  0              0         0    210.0  126.5   87.0  19.16
4237                  0              0         0    269.0  133.5   83.0  21.47

      heartRate  glucose
0         80.0  77.000000
1         95.0  76.000000
2         75.0  70.000000
3         65.0 103.000000
4         85.0  85.000000

```

```

...      ...      ...
4233      66.0      86.000000
4234      65.0      68.000000
4235      84.0      86.000000
4236      86.0      81.966753
4237      80.0      107.000000

```

[4238 rows x 15 columns]

2 Train Test Split

```
[21]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
      ↪2,random_state=42)
```

```
[22]: y_train
```

```

[22]: 3252      0
      3946      0
      1261      0
      2536      0
      4089      0
      ..
      3444      0
      466      0
      3092      0
      3772      0
      860      0
      Name: TenYearCHD, Length: 3390, dtype: int64

```

3 Logistic Regression Algorithm

```
[23]: from sklearn.linear_model import LogisticRegression
      model = LogisticRegression().fit(x_train,y_train)
      model.score(x_train, y_train)
```

```

-----
ValueError                                Traceback (most recent call last)
Cell In[23], line 2
      1 from sklearn.linear_model import LogisticRegression
----> 2 model = LogisticRegression().fit(x_train,y_train)
      3 model.score(x_train, y_train)

File ~\anaconda3\Lib\site-packages\sklearn\base.py:1151, in _fit_context.
      ↪<locals>.decorator.<locals>.wrapper(estimator, *args, **kwargs)
    1144     estimator._validate_params()

```

```

1146 with config_context(
1147     skip_parameter_validation=(
1148         prefer_skip_nested_validation or global_skip_validation
1149     )
1150 ):
-> 1151     return fit_method(estimator, *args, **kwargs)

File ~\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py:1207, in
↳ LogisticRegression.fit(self, X, y, sample_weight)
    1204 else:
    1205     _dtype = [np.float64, np.float32]
-> 1207 X, y = self._validate_data(
    1208     X,
    1209     y,
    1210     accept_sparse="csr",
    1211     dtype=_dtype,
    1212     order="C",
    1213     accept_large_sparse=solver not in ["liblinear", "sag", "saga"],
    1214 )
    1215 check_classification_targets(y)
    1216 self.classes_ = np.unique(y)

File ~\anaconda3\Lib\site-packages\sklearn\base.py:621, in BaseEstimator.
↳ _validate_data(self, X, y, reset, validate_separately, cast_to_ndarray,
↳ **check_params)
    619     y = check_array(y, input_name="y", **check_y_params)
    620     else:
--> 621         X, y = check_X_y(X, y, **check_params)
    622     out = X, y
    624 if not no_val_X and check_params.get("ensure_2d", True):

File ~\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1147, in
↳ check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order, copy,
↳ force_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples,
↳ ensure_min_features, y_numeric, estimator)
    1142     estimator_name = _check_estimator_name(estimator)
    1143     raise ValueError(
    1144         f"{estimator_name} requires y to be passed, but the target y is,
↳ None"
    1145     )
-> 1147 X = check_array(
    1148     X,
    1149     accept_sparse=accept_sparse,
    1150     accept_large_sparse=accept_large_sparse,
    1151     dtype=dtype,
    1152     order=order,
    1153     copy=copy,
    1154     force_all_finite=force_all_finite,

```



```

1155     ensure_2d=ensure_2d,
1156     allow_nd=allow_nd,
1157     ensure_min_samples=ensure_min_samples,
1158     ensure_min_features=ensure_min_features,
1159     estimator=estimator,
1160     input_name="X",
1161 )
1163 y = _check_y(y, multi_output=multi_output, y_numeric=y_numeric,
↪estimator=estimator)
1165 check_consistent_length(X, y)

```

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\validation.py:959, in
↪check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy,
↪force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
↪ensure_min_features, estimator, input_name)
    953         raise ValueError(
    954             "Found array with dim %d. %s expected <= 2."
    955             % (array.ndim, estimator_name)
    956         )
    958     if force_all_finite:
--> 959         _assert_all_finite(
    960             array,
    961             input_name=input_name,
    962             estimator_name=estimator_name,
    963             allow_nan=force_all_finite == "allow-nan",
    964         )
    966     if ensure_min_samples > 0:
    967         n_samples = _num_samples(array)

```

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\validation.py:124, in
↪_assert_all_finite(X, allow_nan, msg_dtype, estimator_name, input_name)
    121     if first_pass_isfinite:
    122         return
--> 124     _assert_all_finite_element_wise(
    125         X,
    126         xp=xp,
    127         allow_nan=allow_nan,
    128         msg_dtype=msg_dtype,
    129         estimator_name=estimator_name,
    130         input_name=input_name,
    131     )

```

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\validation.py:173, in
↪_assert_all_finite_element_wise(X, xp, allow_nan, msg_dtype, estimator_name,
↪input_name)
    156     if estimator_name and input_name == "X" and has_nan_error:
    157         # Improve the error message on how to handle missing values in
    158         # scikit-learn.

```

```

159     msg_err += (
160         f"\n{estimator_name} does not accept missing values"
161         " encoded as NaN natively. For supervised learning, you might
↪want"
        (...)
171         "#estimators-that-handle-nan-values"
172     )
--> 173 raise ValueError(msg_err)

```

ValueError: Input X contains NaN.

LogisticRegression does not accept missing values encoded as NaN natively. For
↪supervised learning, you might want to consider sklearn.ensemble.
↪HistGradientBoostingClassifier and Regressor which accept missing values
↪encoded as NaNs natively. Alternatively, it is possible to preprocess the
↪data, for instance by using an imputer transformer in a pipeline or drop
↪samples with missing values. See [https://scikit-learn.org/stable/modules/](https://scikit-learn.org/stable/modules/impute.html)
↪impute.html You can find a list of all estimators that handle NaN values at
↪the following page: [https://scikit-learn.org/stable/modules/impute.](https://scikit-learn.org/stable/modules/impute.html#estimators-that-handle-nan-values)
↪html#estimators-that-handle-nan-values